

# AWS EC2

## Elastic Compute Cloud

A service that provides secure, resizable compute capacity in the cloud

It is designed to make web-scale cloud computing easier

### EC2 Provides:

- A virtual computing environment
- Preconfigured environments/images called AMI's
  - Allows you to quickly setup a virtual machine
- You can configure CPU, Memory, storage, OS...
- Has secure login via ssh keys, and firewalls to prevent unwanted traffic

The EC2 has a firewall known as Security Groups

- Instance level
- Series of permissions associated with IP's
- Deny by default

## Elastic Block Storage

- Extra storage for an EC2
- Setup and attach individual EBS's

## EC2 AutoScaling

Allows you to scale up or down resources for your application automatically ensuring the availability of your application

### Two ways to autoscale

Vertical: giving the server more resources

- CPU cores / Memory

Horizontal: spin up extra servers, and split the load between them

AWS uses horizontal scaling

## Benefits of Autoscaling

- Improved fault tolerance
- Increased application availability
- Lower costs

## EC2 Autoscaling Groups

Collections of EC2's with similar characteristics

- Use these to increase the number of EC2's
- Use these to decrease the number of EC2's
- Use these to set a fixed amount of EC2's

Launch configurations allow you set the parameters for autoscaling

Scheduled Scaling

Dynamic Scaling

Predictive Scaling